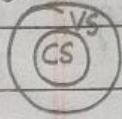


View Serializability

A schedule is view serializable if it is view equivalent to any of the serial schedules.

consistent sch.

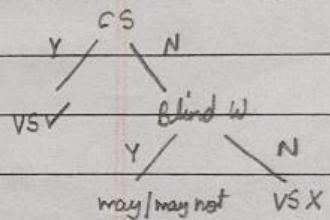


Inconsistent schedules

Conditions for view equivalence:

- Initial reads (for all variables involved) must be the same for both schedules
- Final writes (for all variables involved) must be the same for both schedules
- Updated read / intermediate read: If in schedule S_1 , T_1 is reading value of any variable updated by write operation of T_2 then, in S_2 the order should be same i.e. T_1 must read value of variable after the write operation of T_2 .

- If a schedule is CS, it is VS
- If a schedule is not CS, and doesn't have a blind write, it is not VS. (CS X schedule must have [at least] a blind write [not the only condition to be VS if not VS X])



Example :

S		S'	
T_1	T_2	T_1	T_2
IR R(a)		IR R(a)	
	w(a) ←		w(a) ←
	R(a)		R(b)
	w(a) FW		w(b)
IR R(b)			
	w(b) ←		
	R(b)		R(a)
	w(b) FW		w(a) FW
			R(b)
			w(b) FW

Two transactions ∴ 2! serial schedules are possible

S must be view eq. to any of the 2 serial sch.

IR ✓

FWV

Int. R ✓

S is view equivalent to S'
∴ S is VS

In prev eg
it was CS so
VS as well

Page No _____

Date _____

It is better to check for CS first and then VS.

Eg:

S	T ₁	T ₂	T ₃
IR	R(a)		
	w(a)	w(a) Blind	
		w(a) FW	
			T ₁ → T ₂ → T ₃
			not CS
			contains blind write
			∴ may / may not be VS

Now check for each serial schedule one by one until u get view equivalence

S'	T ₁	T ₂	T ₃	IR
	R(a)			RV
	w(a)			FWV
		w(a)		Int RV
		w(a) _{fw}		(as Int R is not Int)
			w(a) _{fw}	S is v. eq. to S'
				S is VS

Recoverable Schedule

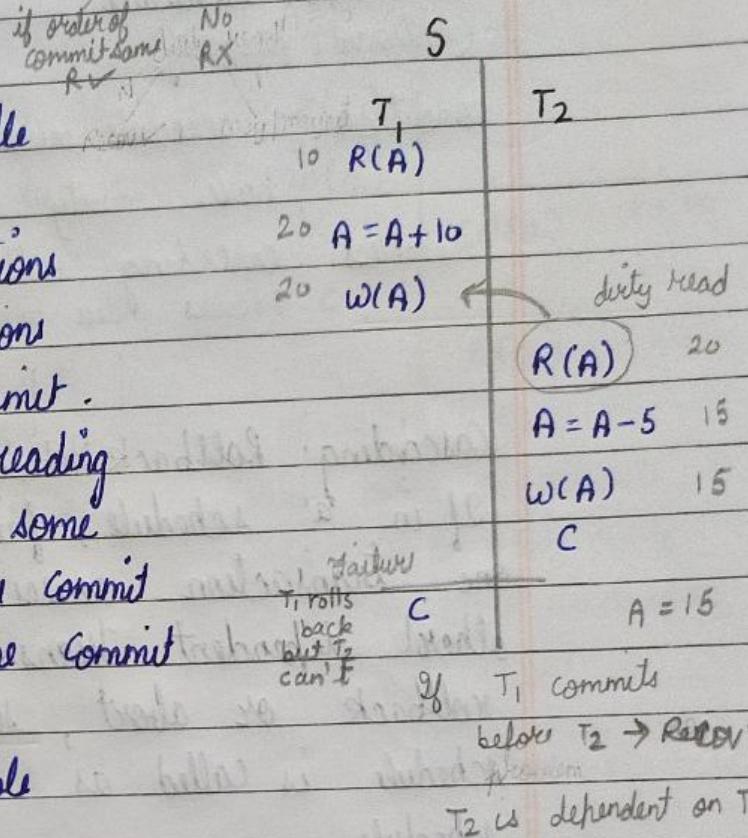
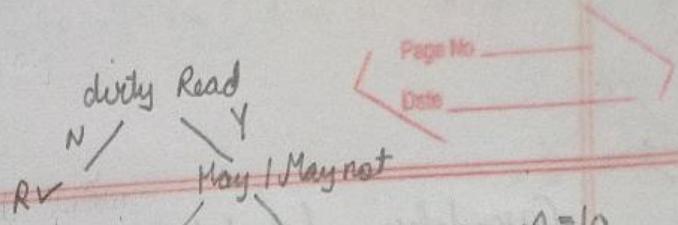
- A schedule which can handle system failures.
- Schedule in which transactions commit only after all transactions whose changes they read commit.
- If some transaction T_j is reading value updated or written by some other transaction T_i , then the commit of T_j must occur after the commit of T_i .

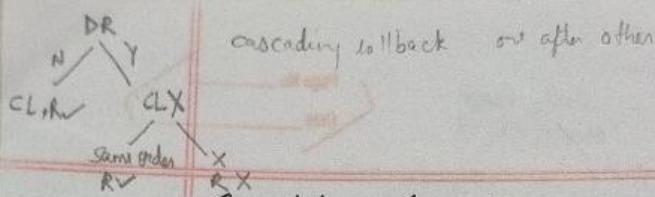
A schedule must be recoverable

It is mandatory property

Condition : must not contain dirty reads

If dirty reads are present \rightarrow Order of rollbacks must be same as order of dirty reads
(cascaded rollbacks)





cascading rollback out after other

Cascadeless Schedule

(optional property)

condition : schedule must not have dirty reads
removal cascading rollbacks occurs due to dirty read)

S

T ₁	T ₂	T ₃
R(A)		
W(A)		
	R(A) ^{DR}	
	W(A)	
		R(A)

Cascading Rollback Schedule :

If in a schedule, failure of one transaction causes several other dependent transactions to rollback or abort, such schedule is called as cascading schedule.

Failure

c

T₁ → T₂ → T₃

Recoverable ✓
Cascadeless ✗

- It leads to wastage of CPU time.
- It occurs due to dirty reads

T ₁	T ₂	T ₃
R(A)		
W(A)		
C		

R(A)^{DR}

W(A)

c R(A)^{DR}

c Recoverable ✓

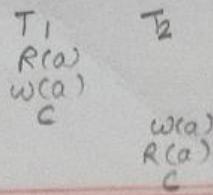
Cascadeless ✓

but concurrency lost

- Dirty reads are eliminated by committing the changes made by a transaction before the other transaction reads the changed value.

- It is an optional property (it reduces concurrency).
- It allows only committed reads. However, it allows uncommitted writes.
- A cascadeless schedule is always recoverable, vice versa not true.

Commit &
read



Page No. _____
Date _____

Strict Schedule

Even if a schedule is cascaded,
it may contain lost updates (due
to blind writes)

A strict schedule is a schedule
in which a write operation
is always followed by a commit.
i.e. ~~another~~ another transaction
performs operations of on a data
value only after changes made
by the ~~earlier~~ first transaction has
been committed.

$a = 10$	T ₁	T ₂
$a = 10$	R(a)	
$a = 15$	w(a)	
		$w(a)$
		$a = 100$
	C	
		R(a) $a = 100$
		C

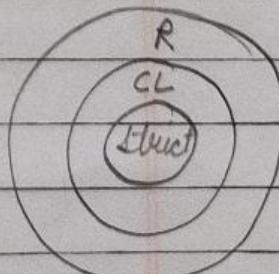
Here when T₁ commits,
it wants to store $a = 15$
but due to blind writes of T₂
 $a = 100$ gets stored

This schedule is cascaded
(no dirty reads)

A strict schedule is a schedule in
which a write changes made by a
write operation of a transaction is
committed before it is read/written
by another transaction.

- To avoid lost updates and dirty
reads.

	T ₁	T ₂
	R(a)	
		R(b)
		w(a)
		w(b)
		C
		R(a)
		C



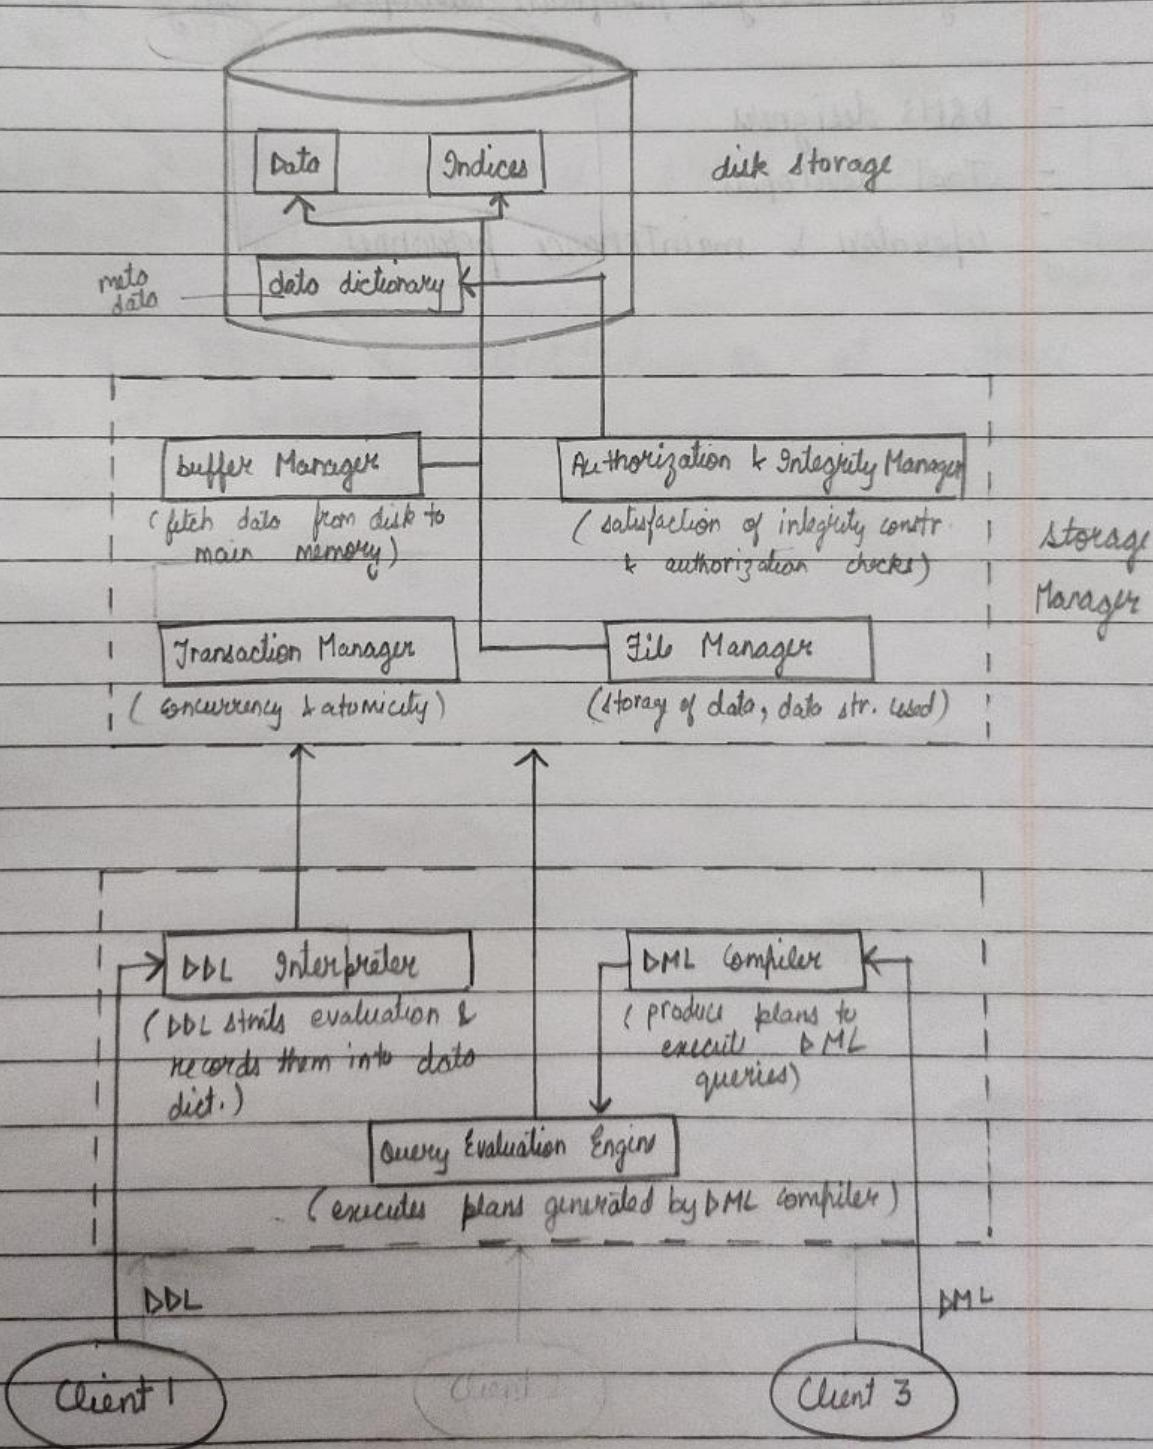
Conclusion :

Inclusion :	
All schedules	
Recoverable	
cascadeless	
strict	
Conflict serializable	
	View serializable

Database system structure

DBMS takes queries, process them & produce plans to implement them, fetches the data from database and sends to appl. prog.

SQL → DML
(non procedural) → compiler → sequence of operations to execute the query (procedural)



People concerned with DBMS

- DBA (Database Administrator) - team
 - manages the resources
 - primary resource - database
 - sec. - DBMS
- Database designers - usually from DBA team
- End users - querying job
- system analysts / software developers - design app program

behind
the
scene
not concerned
with the content

- DBMS designers
- Tool developers
- Operators & maintenance personnel

3 Explain Data abstraction

Data abstraction :

Hiding the complexities of the system from the client & providing an abstract view of only relevant information only is known as data abstraction.

Details of how data is stored & created & maintained is hidden.

Data is DBMS is described at three levels of abstraction :

Name : Payal
Roll no : 23609

Page No _____

Date _____

ASSIGNMENT : I UNIT : I

I Difference between DDL & DML

DDL

- (Data Definition Language)
- It is used to create, delete & modify database schema
- It defines the column (attributes) of the table.
- It doesn't have any further classification
- Commands present in DDL:
Create
Drop
Rename
Alter etc.

DML

- (Data Manipulation Language)
- It is used to insert, update, delete & retrieve data from database.
- It adds or updates the row (tuple) of the table.
- It is further classified as procedural & non-procedural DML
- Commands present in DML:
Update
Insert
Delete
Merge etc.

- It doesn't use where clause.
- Changes are reflected in the schema

- It uses where clause
- Changes are reflected in the instance.

Value set / domain of attribute

Each attribute has its value set V .

If attribute $\rightarrow A$ (simple attribute)

entity set $\rightarrow E$

Value set $\rightarrow V$

$A : E \rightarrow P(V) \rightarrow$ Power set of V

it covers all members of values - null, single, multi

If attribute A is composite

attribute $\rightarrow A$ (composite)

entity set $\rightarrow E$

value sets for its components $\rightarrow V_1, V_2, V_3, \dots, V_n$

value set $\rightarrow V = P(V_1) \times P(V_2) \times P(V_3) \times \dots \times P(V_n)$

$A : E \rightarrow P(V) = P[P(V_1) \times P(V_2) \times P(V_3) \times \dots \times P(V_n)]$

B

KB 2^5
 MB 2^{10}
 GB 2^{20}
 TB 2^{30}
 PB 2^{40}

ER Diagram

Conversion of ER \rightarrow relational

Basics of rel. model & FD

designing

Keys / types

Normalisation

Loskau decomposition & dependency preserving

Indexing & physical str.

SQL, RA, RC

Transaction

Concurrency Control

OLAP data warehouse
legacy data

(Online Analytical Processing)

- historical data
- subject oriented
- decision making
- size TB, PB
- CEO, Managing Director
- General Manager - access
- only read operations

OLTP database

(Online Transaction Processing)

- current data
- application oriented
- day to day operations
- MB, GiB
- clerk, managers
- read / write

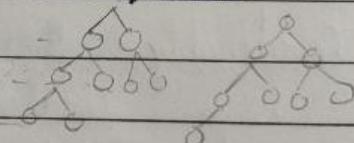
Complete binary tree - all full except last level - as left as possible

$$\text{for } h \text{ ht} \quad \text{max nodes} = 2^{h+1} - 1$$

$$\text{min nodes} = 1, 2, 4, 8 = 2^h$$

$$\text{for } n \text{ nodes} \quad \text{min ht} = \lceil \log_2(n+1) \rceil - 1$$

$$\text{max ht} = \log_2 n$$

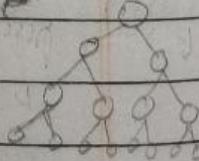


- Full / Proper / Strict - each node has 2 nodes except leaf \leftarrow either 0 or 2 children

$$\Rightarrow \text{no. of leaf nodes} = \text{no. of internal nodes} + 1$$

$$\text{for } h \text{ ht} \quad \text{max nodes} = 2^{h+1} - 1$$

$$\text{min nodes} = 1, 3, 5, 7, \dots = 2h + 1$$



$$\text{for } n \text{ nodes} \quad \text{max ht} = (n-1)/2$$

$$\text{min ht} = \lceil \log_2(n+1) \rceil - 1$$

Perfect BT - 2 child left at same level

Skewed BT - all has 1 child

Left skewed BT

CBT ✓

WB tree

FBT ✓

Right skewed

NORMALISATION

Process of removing redundant data from the relations to improve storage efficiency, data integrity, and scalability and speed.

It usually involves decomposition of relations.

First Normal Form

→ A table is said to be in first normal form when it follows following conditions :

- Values of all the attributes must be atomic.
- Every column must have unique name.
- Every record of a relation must be unique.
- All values of a particular column must be from same domain.
- Order of rows & columns is irrelevant.
- A relation must have a primary key.

→ When an ER model is mapped into relational model, it is present in 1NF.

→ 1NF can be attained by

- Creating multiple columns for multi-valued attributes and composite attributes.
- Creating multiple tables (for multivalued attributes)

Example:

NOT IN 1NF

→	Roll	Name	Contact	
initially	101	A	9562338912, 9123457891	not in 1NF
but	102	B	9143287321, 8215331448	Contact is multi-valued

↓

→	Roll	Name	Contact 1	Contact 2	
	101	A	9562338912	9123457891	
	102	B	9143287321	8215331448	

→	SId	Name	Hobby	
	1	X	Singing, Painting, Reading	
	2	Y	Dancing, Writing	

SId	Name	S-No	Hobby	SId
1	X	101	Singing	1
2	Y	102	Painting	1
		103	Reading	1
		104	Dancing	2
		105	Writing	2

→	Roll	Name	Course	
	1	A	CN/OS	no problem
	2	B	DBMS/DS	

↓

	Roll	Name	Course	
	1	A	CN	problem
	2	A	OS	problem
	3	B	DBMS	problem
	4	B	DS	

Second Normal Form

- A relation is said to be in 2NF if
 - It is in 1NF
 - It doesn't contain partial dependencies
- Partial dependency: A functional dependency is said to be a partial dependency if a part of candidate key determines the non-prime attributes.
- Problem with partial dependency

	A	B	X
1	Null	a	
2	2	b	
3	Null	c	
4	5	d	

Here AB combined is a composite CK & can determine X. But part of AB (i.e. 'A' or 'B') can not determine X because they contain 'Null' values.
- To find whether a table is in 2NF
 - (when it is given that table is in 1NF)
 - Find all the CKs
 - Find prime and non prime attributes
 - Non prime attributes must not lie involved in partial dependencies.

Example : R(ABC) $A \rightarrow C$

- Candidate key : $(AB)^+ = ABC$
 $\therefore AB$ is a CK

A and B are essential attributes of CK \therefore it is the only CK of given relation.

- Prime attributes : A, B
 - Non prime attributes : C

Since, in FD $B \rightarrow C$, B is a part of CK and C is a non prime attribute, it is a partial dependency. \therefore given relation is not in 2NF.

\rightarrow To convert in 2NF

$$R(ABC) \Rightarrow R_1(B.C) \text{ & } R_2(A.B)$$

$$B \rightarrow C \qquad \qquad B \rightarrow C$$

A	B	C		A	B		B	C
a	1	x		a	1	x	1	x
b	2	y		b	2	y	2	y
a	3	z		a	3	z	3	z
c	3	z		c	3			R_2
d	3	z		d	3	z	z	
e	3	z		e	3			

Here, R contains redundant data in the records corresponding to $B = 3$. After discarding it,

After decomposition, redundancy is reduced and 2NF is attained.

Third Normal Form

A relation is said to be in 3NF if,

- It is in 2NF
- It does not contain transitive dependencies

Transitive dependency: A functional dependency is said to be transitive if a non prime attribute determines another non prime attribute.

Problem with transitive dependency

A	B	C
a	1	2
b	1	2
c	1	2
d	2	4
e	2	4

Here A is the CK and B, C are non prime attributes where $B \rightarrow C$. It leads to redundancy of data.

To find whether a relation is in 3NF

(when already present in 2NF)

- Find all the CKs
- Find prime and non prime attributes
- Check whether any non prime attribute is being determined by another non prime attribute

or

Check for every dependency from $\alpha \rightarrow \beta$

- either α is a super key
- or β is a prime attribute

example:

$R(ABC)$ $A \rightarrow B$ $B \rightarrow C$ in relation A

Candidate key: $A^+ = ABC$ is in 3NF

A is a CK (candidate key) to make 3NF

since A is an essential attribute $\therefore A$ is the only CK.

prime attributes: A is prime attribute

Non Prime attributes: B, C

In F_B $B \rightarrow C$ B which is a non prime attribute determines AC which is another non prime attribute \therefore it is a transitive dependency
Hence given relation is not in 3NF

To convert into 3NF

$R(ABC) \rightarrow R_1(BC), R_2(AB)$

$A \rightarrow B$

$B \rightarrow C$

$B \rightarrow C$

$A \rightarrow B$

R_1 is relation 1 relation 1 of

(this is primary, generate relation)

A	B	C	A	B	C	B	C
a	1	x	a	1	x	1	x
b	1	x	b	1	x	2	y
c	1	x	c	1	x	3	z
d	2	y	d	2			
e	2	y	e	2			
f	3	z	f	3			
g	3	z	g	3			

R

R_1 R_2

BCNF (Boyce Codd Normal Form)

day

PM
PM

→ A relation is said to be in BCNF if

- It is in 3NF
- Prime attributes are not determined by any non prime attribute or part of CK.

→ To determine whether a relation is in BCNF
(when already present in 3NF)

- Find all the CKs
- Find prime & non prime attributes.
- Check whether any prime attribute is being determined by non prime or part of CK.

Example: R(ABC) $AB \rightarrow C$ $C \rightarrow B$

$$\text{Candidate keys: } (AB)^+ = ABC \\ (AC)^+ = ABC$$

∴ AB & AC both are CKs.

Prime attributes : A, B, C

Non prime attributes : -

Here in R, $C \rightarrow B$, C is a part of CK and determines a prime attribute B
∴ It is not in BCNF

or In $F \triangleright \alpha \rightarrow \beta$
 α should be super key.

Authenti Autho access control DAC MAC RBAC Models
invasion detection SQL injection

DATABASE SECURITY

Database security refers to the range of tools, controls and measures designed to establish and preserve database confidentiality, integrity and availability.

It handles various issues such as

- legal and ethical issues
- Policy issues
- System - related issues
- The need to identify multiple security levels

Threats to databases :

#

Loss of integrity : Database integrity refers to the requirement that information be protected from unauthorised modification.

Use of corrupted data could result in inaccuracy, fraud or erroneous decisions.

#

Loss of availability : Database availability refers to making objects available to a human user or a program to which they have a legitimate right.

#

Loss of confidentiality : Database confidentiality refers to the protection of data from unauthorised disclosure.

Control Measures:

- Access control
- Inference control
- Flow control
- Data encryption

ACCESS CONTROL

Access control is a method of allowing access to sensitive data of database only to those people who are allowed to access the data and to restrict the access to unauthorised persons.

Authentication : It is a method of verifying the identity of a person who is accessing the database.

- It usually requires user's login details.
- It verifies whether the user is authenticated or not.
- eg: log in id , password

It can be
→ single factor authentication

(only user id and password is required)

→ 2 factor authentication

(more info like OTP etc is required)

Authorisation : It is a method of verifying the authorities of a user for accessing the resources.

- It requires user's privilege or security levels.
- It determines the permissions user have.

Access control

- Whenever a user needs to access a database system, he has to apply for a user account.
- DBA creates the user account and password if there is legitimate need to access the data.
- The user must login with the account number and password which is validated by the DBMS.

audit trail

It is a record which contains all the updates applied to the database and particular user who applied each update.

It is useful for database audit or in case of system failures.

Access control Mechanisms :

Discretionary Access Control (DAC)

It is a type of security access control based on granting and revoking of privileges.

Types :

- > Account level: At this level, the DBA specifies the particular privileges that each account holds independent of relations in the database.

- It may include CREATE SCHEMA, CREATE TABLE, CREATE VIEW, DROP, MODIFY privilege etc

2) Relation level: At this level, DBA can control the privilege to access each individual relation or view in the database.

- It may include SELECT on R, Modification on R etc

• Grant: It is used to provide a user permission to use a privilege

• Revoke: It is used to take back the permission granted

Propagation of privilege

A privilege on R can be given with or without grant option. If a privilege is provided with a grant option, the privilege can be propagated.

But if the privilege is revoked later, it must be taken back from all the accounts to which it was propagated.

Mandatory Access Control (MAC)

- It is a hierarchical model. The hierarchy is based on security levels
- All users are assigned a security or clearance level and all objects are assigned a security label.
- Generally four security classification levels are used:
 - Top secret (TS)
 - Secret (S)
 - Confidential (C)
 - Unclassified (U)

$$TS \geq S \geq C \geq U$$

Simple security property

- A subject s is not allowed to read access to an object o unless
 $\text{Class}(s) \geq \text{Class}(o)$
- This implies that no subject can read an object whose security classification is higher than the subject's security clearance.

Star property

- A subject s is not allowed to write an object o unless
 $\text{class}(s) \leq \text{class}(o)$
- It prohibits a subject from writing an object at a lower security classification than the

subject's security clearance

- Violation to this rule would allow information to flow from higher to lower classifications

MAC ensures a high degree of protection

- they are suitable for military and high security types of applications

Role-Based Access Control (RBAC)

- Role based access control restricts the access based on a person's role within an organization rather than individual users.
- Access can be based on several factors such as authority, responsibility, job competency etc.
- Various roles are created and security privileges are granted to the role name and any individual assigned to this role would automatically have those privileges granted
- Users may also be assigned to multiple groups roles.
- Separation of duties can be implemented by providing mutually exclusive roles.
 - Mutually exclusive roles can be:
 - authorisation time exclusion (static)
roles that are mutually exclusive can't be part of a user's authorisation at the same time under authorisation time exclusion.
 - runtime exclusion (dynamic)
roles can be authorized to one user but can't be activated at same time.

Name : Payal
Class : CSE 2 (Sem. III)
Roll No : 23609

UNIT : 1

Ques: Explain ER diagram. Define various types of attributes and relationships.

UNIT : 2

Ques: Define Normalisation. Explain various normal forms.

UNIT : 3

Ques: Define Transaction. Explain ACID properties
Also explain different states of a transaction.

UNIT : 4

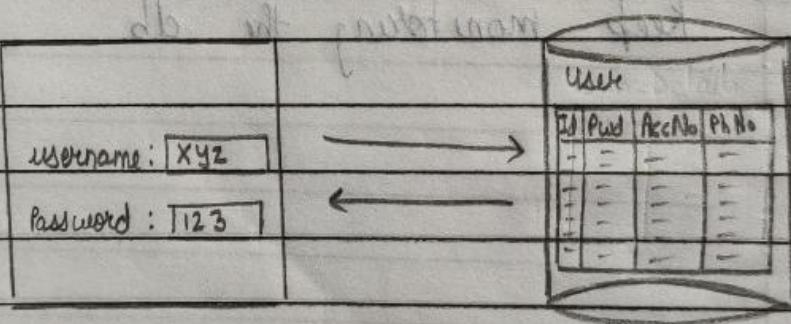
Ques: Explain Database security? Briefly explain DAC, MAC & RBAC models.

SQL Injection

SQL injection is a type of attack on database in which user's data is exploited by user through web page inputs by injecting malicious SQL commands into it.

- Web programs and applications that access a database can send commands and data to databases as well as display data retrieved from the database through the web browser.
- In SQL injection, a string input is injected through the application which manipulates the SQL statement to attacker's advantage.

Example :



Now, if the attacker enters:

username : " OR 1 = 1 --
to close
username
string

A condition which
is always true
is injected with
the OR operator
so that the result
is always true.

earlier query : select * from users where Id = "xyz"
and ~~pass~~ pwd = "123"

later query : select * from users where Id = "xyz"
or 1 = 1 -- and pwd = "123"

always true ∴ attacker will login successfully even if Id & pwd are incorrect

Risks :

mitigation 100

- Database fingerprinting
- Denial of service
- Bypass authentication
- Identifying injectable parameters (gathers info abt back end)
- Executing remote commands
- Performing privilege escalation

Prevention

- Bind variables
- tools & software → security event manager
NetSparker Vulnerability scanner
- limit the size of fields
- Keep monitoring the db

Intrusion Detection

An Intrusion Detection System (IDS) is a system that monitors network traffic for suspicious activity and issues alerts when such activity is discovered. It is a software application that scans a network or a system for harmful activity or policy breaching.

Any malicious venture

B TREE (Balanced Tree)

- It is a self balancing search tree.
- It is used to maintain multi-level indices for the data stored in secondary memory.
- It reduces the disk access time.
- Size of a node is kept equal to block size of disk.
- Height of B-tree is low. (It is broader)

Properties :

- All leaf nodes are at same level.
- It is defined by order 'm' which depends on the block size of disk.
- Maintains data in sorted order.
- If order of B-tree = m
 - min children : leaf = 0
root = 2
 - max children = m
 - min keys : root = 1
all other $\lceil \frac{m}{2} \rceil - 1$
 - max keys = $m - 1$
- Time complexity for search, insert, delete is $O(\log n)$ $n = \text{nodes}$

Each node contains search key as well as record pointer
& link to child nodes.

structure of node

$m=4$

BP	key	DP	BP	key	DP	BP	key	DP	BP
----	-----	----	----	-----	----	----	-----	----	----

BP = Block Pointer / Tree pointer
points to the child node

key = data value used for searching

DP = Data pointer / Record pointer
points to the record where
actual data is stored

Ques: Consider a B-Tree with key size = 10 bytes,
block size 512 bytes, data pointer is of size 8
bytes and block pointer is 5 bytes. Find the
order of B-Tree

node size \leq block size

size of block = 512

$$\begin{aligned}\text{node size} &= n(\text{size of BP}) + (n-1)[\text{size of key} + \text{size of DP}] \\ &= n(5) + (n-1)[10 + 8] \\ &= 5n + 18n - 18 = 23n - 18\end{aligned}$$

Here n = order of B-tree

$$512 \geq 23n - 18$$

$$n \leq 530/23$$

$$n \leq 23.04$$

$\therefore n$ must be 23

only leaf node contains record pointers.

B⁺ TREE

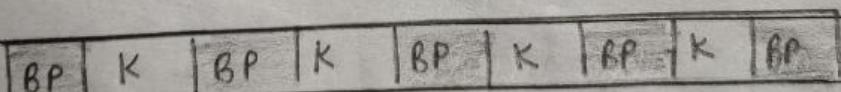
In B tree data is stored in leaf node as data record / data pointers are stored along with the keys in leaf nodes as well as internal nodes. This reduces the no. of entries of a node ∵ increases height

- But in B⁺ trees data pointers are stored only in leaf node.
- Leaf nodes contain all the key values along with their corresponding record pointers.
- So leaf nodes form the first level index and internal nodes form other multi levels.
- Height remains balanced & less as compared to B Tree
- All leaf nodes are linked.

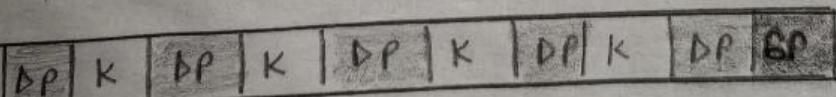
Structure of nodes :

for internal node :

$$m=5$$



leaf node :



search access time is lesser

CONCURRENCY CONTROL TECHNIQUE

The concurrency control protocols ensure the atomicity, consistency, isolation durability & serializability of concurrent schedules.

concurrency ✓ properties of time ↓ logic ✓

- Time stamp protocol
- Lock based protocol
 - 2PL
 - graph based
- Validation protocol

TIME STAMP PROTOCOL

- Time stamp ordering protocol orders the transactions in the increasing order of their creation i.e. based on their time stamps.
- The priority of older transaction is higher that's why it executes first in case of conflicts.
- System clock is used for timestamps bcz it is unique & never repeats itself.

→ Two kinds of time stamps are maintained:

1) Time stamp with transaction

- for each transaction T_i , a timestamp is associated denoted by $TS(T_i)$ based on the system clock.

- this timestamp remains constant throughout the execution of that transaction.

- if T_j enters after T_i then,

$$TS(T_i) < TS(T_j)$$

- it helps in determining the order of serializability

2) Time stamp with data item

For each data item α , 2 timestamp are maintained.

- w-timestamp (α): The largest timestamp of any transaction that executed $\text{write}(\alpha)$.

- r-timestamp (α): The largest timestamp of any transaction that executed $\text{read}(\alpha)$.

Rules

If T_i request for $\text{read}(\alpha)$ Two cases:

1) if $TS(T_i) < WTS(\alpha)$

means T_i wants to read a value already overwritten by a transaction having lesser priority

In this case Request must be rejected & T_i must rollback.

$$\text{eg: } TS(T_i) = 5$$

$$WTS(\alpha) = 10$$

$$TS(T_\alpha) = 10$$

$$T_\alpha \text{ is}$$

w(α)

not allowed

In case of conflicts T_i with higher priority i.e. lesser TS

must perform the operation

2) if $TS(T_i) \geq WTS(\alpha)$ if equal \rightarrow same transaction

means T_i wants to read a value already overwritten by a transaction having higher priority

In this case operation is allowed & $R TS(\alpha)$ will be updated to $\max(TS(T_i), R TS(\alpha))$

$$\text{eg: } TS(T_i) = 10$$

$$WTS(\alpha) = 5$$

$$TS(T_\alpha) = 5$$

$$TS(T_i) = 10$$

$$T_\alpha$$

$$T_i$$

$$w(\alpha)$$

$$R(\alpha)$$

allowed

If T_i request for write (α): 4 cases

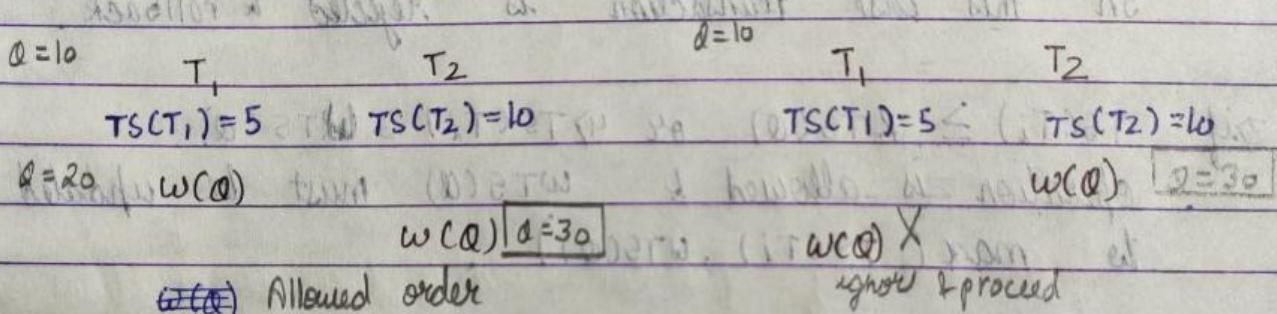
- 1) if $TS(T_i) < RTS(\alpha)$
means value of α is already read by another transaction with lower priority.
In this case transaction is rejected & rollback.
- 2) if $TS(T_i) < WTS(\alpha)$
means value of α is already overwritten by a transaction having lesser priority.
In this case transaction is rejected & rollback.
- 3) if $TS(T_i) \geq RTS(\alpha)$ or $TS(T_i) \geq WTS(\alpha)$
operation is allowed & $WTS(\alpha)$ must be updated
to $\max(TS(T_i), WTS(\alpha))$.

Properties

- It ensures conflict serializability.
- It ensures view serializability.
- Possibility of dirty read, no restriction on commit.
∴ irrecoverable and cascading rollback schedules are possible.
- Either an operation is allowed or rejected, so deadlock doesn't occur.
- Every Timestamp protocol-schedule is CS but all CS schedules may not be generated by TS.
- May lead to starvation & relatively slow.

THOMAS WRITE RULE

- It modifies time-stamping protocol to make some improvements & may generate schedules which are VS but not CS & provides better concurrency.
- Modifies time-stamping protocol in obsolete write case:
 T_i requests for write (α)
 if $TS(T_i) < WTS(\alpha)$
- In this case rather than rollback of transaction, write operation is ignored & transaction is proceeded.



Assume $T_1 \rightarrow w(\alpha) \rightarrow \alpha = 20$
 $T_2 \rightarrow w(\alpha) \rightarrow \alpha = 30$

$$TS(T_1) = 1 \quad TS(T_2) = 2 \quad TS(T_3) = 3$$

e.g.: $T_1 \quad T_2 \quad T_3$

$R(\alpha)$

$w(\alpha)$

$Xw(\alpha)$
ignore

$w(\alpha)$

$T_1 \rightarrow T_2$

$T_2 \rightarrow T_3$

$T_1 \rightarrow T_3$

$GSS X$

20 like T_1 and T_2 w. T_3 with $\alpha = 10$ IR ✓
 $R(\alpha)$ 2+ not belations $\neq R$ ✓
 $w(\alpha)$ w. T_1 and T_2 Inter R X
 $w(\alpha)$ w. T_1 and T_2 VS ✓
 $w(\alpha)$

LOCK BASED PROTOCOL

In lock based protocol, a transaction applies lock on the data item before performing any operation on it & then unlocks it after completion of the operations.

Types of locks:

	shared	exclusive
shared	T	F
exclusive	F	F

1) Shared lock Lock - S(A)

(Read-only lock)

- In this lock, data items can only be read.
- It can be shared among the transactions i.e. if a transaction applies shared lock on any data item, other transactions may also apply the shared lock on the same data item.

2) Exclusive lock Lock - X(A)

- In this lock, data items can be read as well as written.

- This lock can't be applied by multiple transactions on same data item. Only one transaction can apply exclusive lock on any data item at a time.

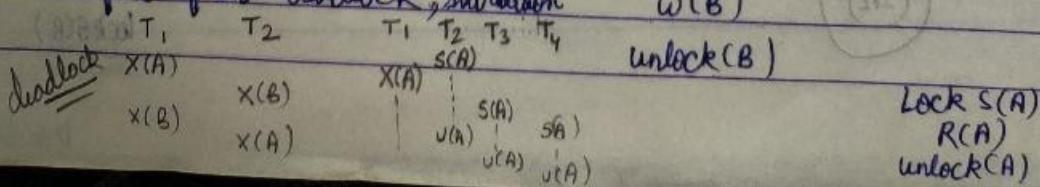
example

- It doesn't ensure CS, VS, recoverability or cascadeless rollback.

- If data item is unlocked before committing, it may lead to inconsistency.

If data item is unlocked only after committing, concurrency is affected.

- May not be free from deadlock, starvation



Due to these disadv., transaction must follow some set of rules for locking & unlocking data item eg 2PL, graph based.

2-PHASE LOCKING / Basic 2PL

In 2 phase locking protocol, each transaction in a schedule has 2 phases:

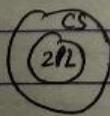
1) growing phase: In this phase, transaction can only obtain locks but can't release any lock.

2) shrinking phase: In this phase, transaction can only release locks but can't obtain any lock.

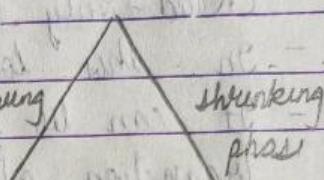
- So transaction first enters into the growing phase, then obtains all the locks, then reaches the lock point & finally moves to shrinking phase & releases all the acquired locks
- It can perform read, write operation in both the phases.

Properties:

- Ensures CS & VS
- Order of serializability is the order in which transactions reaches lock point.
- May generate unrecoverable schedules or cascading rollbacks.
- Do not ensure freedom from deadlock/starvation
- Every 2-phase locking - protocol - schedule is CS but every CSS may not be generated from 2-Phase locking protol.



(a) 1
(a) 2
(a) 3
(a) 4



T1
lock S(A)
lock X(B)
unlock X(A)
unlock (B)

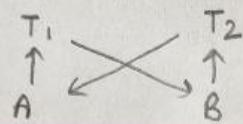
T2
lock R(S(A))
lock X(D)
unlock (A)
unlock (D)

lock X(B)
R(A)
lock S(B)

lock X(B)
R(B)
W(B)

lock S(B)
waiting
lock X(A)
waiting
Deadlock

deadlock occurs when



Conservative / static 2PL resolves deadlock problem

- In this protocol, there is no growing phase, transaction first acquires all the locks & then starts from the lock point and shrinking phase continues as usual.
- If all locks required by a transaction are not available, transaction must release all the locks acquired & wait.
- Prior knowledge of all the data items required is necessary

- It ensures CS, VS
- Independent from deadlock
- Possibility of irrecoverable schedules & cascading rollbacks.

Rigorous 2PL resolves unrecoverability & cascading rollbacks

In this protocol, there is no shrinking phase, transaction holds the lock until it is committed.

- ensures CS & VS
- ensures recoverability & cascadeless rollbacks.
- suffer from deadlock
- inefficient

	T1	T2
X(A)		
R(A)		
W(A)		
C		
U(A)		
S(A)		
R(A)		
U(A)		

Strict 2PL

improvement over Rigorous 2PL

- In this protocol, partial shrinking phase is present.
- In shrinking phase unlocking of exclusive locks are not allowed but unlocking of shared locks can be done as they provide only read permission which do not update the data item hence dirty read doesn't occur.

All properties are same as Rigorous 2PL except that it is efficient

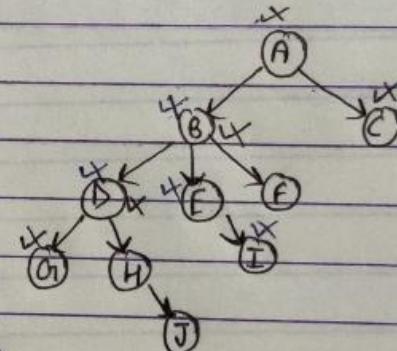
GRAPH BASED PROTOCOL

- Graph based protocols require prior knowledge about the transactions.
- The order in which the database items α will be accessed must be known.
- Partial ordering is imposed on the data items (\rightarrow)
 $\triangleright \sim \{d_1, d_2, \dots, d_n\}$
 if $d_i \rightarrow d_j$, then any transaction accessing both d_i & d_j must access d_i before d_j .
- Partial ordering may be applied through logical or physical organisation (e.g. linked list, graph).
- After partial ordering, set of all data items Δ will be viewed as directed acyclic graph called database graph.
- Assumptions:
 graph must be rooted tree
 only exclusive lock is considered

Tree Protocol

Following rules must be applied

- First lock by T_i may be applied on any data item.
- Subsequently, a data item α can only be locked by T_i if the parent of α is currently locked by T_i .
- Data item may be unlocked at any time.
- Data item α that has been locked & unlocked by T_i can't be relocked by T_i .



T_1 T_2

lockX(B) lockX(A)

lockX(D) lockX(B)

lockX(E) lockX(C)

unlock(B) unlock(A)

lockX(I) lockX(D)

unlock(D) unlock(B)

unlock(F) lockX(G)

unlock(I) unlock(D)

unlock(G) unlock(G)

unlock(C) unlock(C)

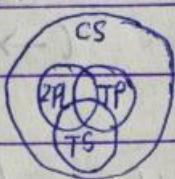
- CS & VS ✓

- freedom from deadlock

- recoverability & cascadelock not ensured.

- data item is unlocked as soon as it has been used \rightarrow waiting time \downarrow & concurrency \uparrow

- A transaction may have to lock data items that it doesn't access to access the desired data item.
- leads to overhead waiting time & decrease in concurrency.
- Transaction must know exactly what data items are to be accessed.



- Recoverability & cascadelessness can be attained by not unlocking before commit.

Question :

Time Stamp Protocol

100	200	300		A	B	C
T ₁	T ₂	T ₃	RTS	∅ ¹⁰⁰	∅ ²⁰⁰ ³⁰⁰	∅ ¹⁰⁰
R(A) ^{100>0}			WTS	∅ ³⁰⁰	0	∅ ¹⁰⁰
	R(B) ^{200>0}				↓	
W(C) ^{100>0} ^{200>0}		R(B) ^{300>0}	RTS	A	B	C
R(C) ¹⁰⁰⁼¹⁰⁰	X		WTS	100	300	100
	w(B) ^{200<300}			300	0	100
		w(A) ^{300>100} ^{200>0}				

T₁ & T₃ are executed successfully & T₂ will rollback

Multi-Version Concurrency Control

- In this technique, multiple versions of a data item are maintained and values of data item are maintained when the item is updated.
- When transaction requires access to an item, an appropriate version is chosen to maintain the serializability of the schedule.
- Due to this some read operations, that would be rejected in other techniques, can still be accepted by reading an older version of item to maintain serializability.
- Multi-version technique requires more storage.
- But older versions may have to be maintained for recovery etc.

Multi-version Technique based on timestamp ordering

In this method, several versions of each data item is maintained & for each version a_i following data is recorded:

- Value of version a_i
- Read TS(a_i) - largest TS of trans. that read a_i
- Write TS(a_i) - largest TS _____ wrote a_i

Rules:

If T_i issues $\text{Read}(a_i)$ request

For $\text{Read}(a_i)$, the recent version of a is provided
if $\text{TS}(a_i) \geq \text{RTS}(a_i)$

Database Recovery

Database systems are subject to failures

- 1 Failures include

- system failure

- transaction errors

- viruses

- catastrophic failure

- power failure etc.

A database system must be capable of recovering from these failures ensuring the transaction properties like atomicity, durability etc.

Recovery techniques are dependent on system log

System log: It is a special file which contains all the information about a transaction - start of transaction, end of transaction & all the operations performed by the transaction.

- This information is important for recovery.

System log contains following entries:

- \langle Transaction T, start \rangle : This entry records that transaction T has started the execution.

- \langle Transaction T, new data item X, new value \rangle : This entry depicts that transaction T has overwritten the value of X to the new value.

- \langle Transaction T, data item X, old_value, new_value \rangle : This entry depicts that transaction T has changed value of X from old_value to new_value.
old_value - image of X
new_value - afterimage of X

- $\langle \text{Transaction } T, \text{ commit} \rangle$: This entry depicts that transaction T has completed all the operations & committed successfully.

- $\langle \text{Transaction } T, \text{ abort} \rangle$: abort

- **Checkpoint**: At this point all the previous logs are removed from the system & stored permanently on storage disk.

Undoing: If transactions are undone, log entry $\langle \text{Transaction } T, \text{ data-item } x, \text{ old-value}, \text{ new-value} \rangle$ is used & value of data item x is set to old-value.

redo: If transactions are redone, log entry $\langle \text{Transaction } T, \text{ data-item } x, \text{ new-value} \rangle$ or $\langle \text{Transaction } T, \text{ data-item } x, \text{ old-value}, \text{ new-value} \rangle$ is used & value of x is set to new-value.

Deferred Database Modification

- It is a log based recovery technique.
- In deferred update technique, database on disk is not updated until commit point.
- if failure occurs before commit, changes are not made into the database \therefore UNDO is not required.
- But if failure occurs after commit, REDO is required.

Example:

T_1	$A = 100$	$B = 200$	log file
$R(A)$			$\langle T_1, \text{ start} \rangle$
$A = A + 100$			$\langle T_1, A, 200 \rangle$
$W(A)$	200		$\langle T_1, B, 400 \rangle$
$R(B)$	200		$\langle T_1, \text{ commit} \rangle$
$B = B + 200$		400	
$W(B)$	400		
Commit			
			after commit

Case I Failure occurs before commit - no action

Case II Failure occurs after commit - redo

$A = 200$ & $B = 400$ are stored

Immediate Database Modification

It is a log based recovery technique

In immediate update technique, changes are made permanent onto the disk as soon as operation is executed.

If transaction fails before commit, undo is required.

If transaction fails after commit, redo is required.

Example:

T_1

$R(A)$

$A = A + 100$

$w(A)$

$R(B)$

$B = B + 200$

$w(B)$

$w(B)$

Commit

$A = 100$

$B = 200$

$A = 200$

$B = 200$

$A = 200$

$B = 400$

log file

$\langle T_1, \text{start} \rangle$

$\langle T_1, A, 100, 200 \rangle$

$\langle T_1, B, 200, 400 \rangle$

$\langle T_1, \text{commit} \rangle$

Case I

Failure occurs before commit - undo

$A = 100$ & $B = 200$ are stored.

Case II

Failure occurs after commit - redo

$A = 200$ & $B = 400$ are stored.

$\langle \text{trans}, T \rangle$

$\langle \text{end}, A, T \rangle$

$\langle \text{end}, B, T \rangle$

$\langle \text{trans}, T \rangle$

(A) 100

add A = A

(A) 200

(B) 400

add B = B

(B) 400