

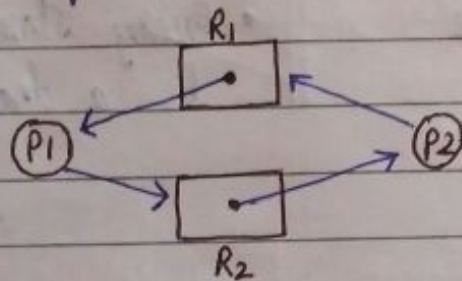
## DEADLOCK

resource requested is held by some waiting process ....

It is a situation where a set of processes are blocked because each process is holding a resource and waiting for another resource acquired by some other waiting process.

- waiting time is infinite and no progress at all (for all the processes present in deadlock)

example



## Coffman Conditions

(necessary conditions for deadlock to occur) <sup>not sufficient</sup>

- 1) Mutual Exclusion At least one resource must be non-sharable.
- 2) Hold & Wait A process can hold a resource & waits for the other resources held by some waiting process.  
(applies to all processes in the deadlock)
- 3) No preemption Resources can't be taken from a process unless it releases the resource.
- 4) Circular wait Set of processes waits for each other in a circular form.

### Deadlock

- no process proceeds
- infinite waiting
- All the resources not in deadlock are held by waiting processes & are not being used.
- All deadlocks are starvation.

### Starvation

- high priority processes proceed while LPP are blocked
- long time waiting, not
- Resources are used being used continuously by high priority processes.
- Every starvation need not be a deadlock



## Resource Allocation Graph

It is a set of

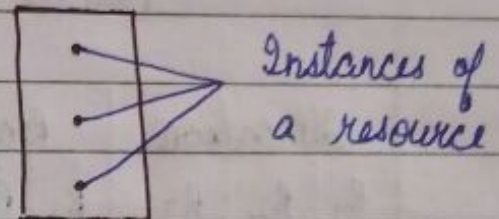
→ Vertices  $\begin{cases} \rightarrow \text{Processes } \{P_1, P_2, P_3, \dots, P_n\} \\ \rightarrow \text{Resources } \{R_1, R_2, R_3, \dots, R_n\} \end{cases}$

→ Edges  $\begin{cases} \rightarrow P_i \rightarrow R_j & (\text{Request edge}) \\ \rightarrow R_j \rightarrow P_i & (\text{Assignment edge}) \end{cases}$

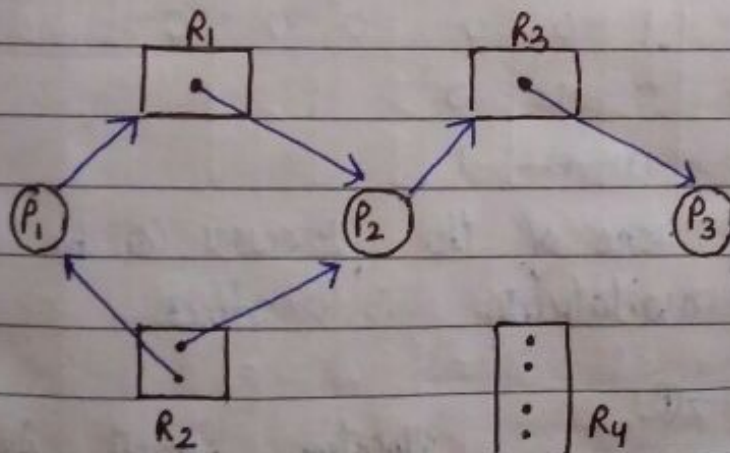
Symbols:

$R_i$  Resource

$P_i$  Process



Example:



Here,  $P = \{P_1, P_2, P_3\}$

$R = \{R_1, R_2, R_3, R_4\}$

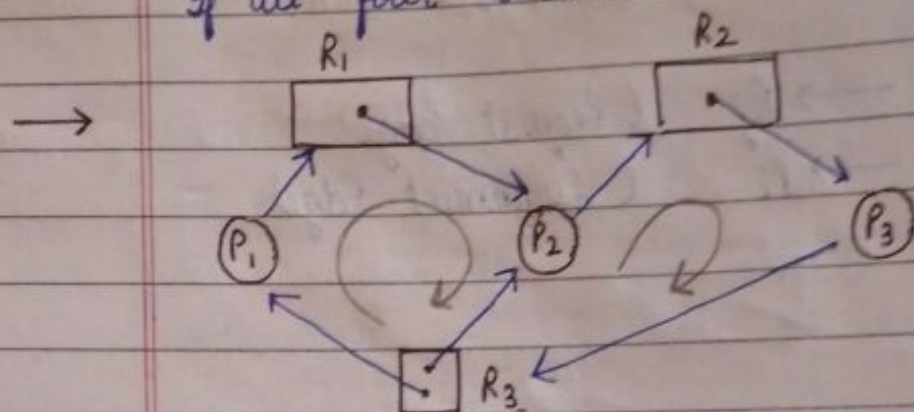
$E = \{P_1 \rightarrow R_1, P_2 \rightarrow R_3, R_1 \rightarrow P_2, R_2 \rightarrow P_1, R_2 \rightarrow P_2, R_3 \rightarrow P_3\}$

Request edge

Assignment edge

## Detection of Deadlock

If any of the four Coffman conditions is not fulfilled, deadlock doesn't occur.  
If all four conditions hold true, deadlock MAY occur.

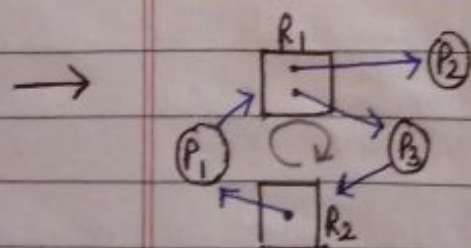


First check for cycle.

If cycle exists, check whether any process has progress.

	Allocation			Request			Availability		
	R <sub>1</sub>	R <sub>2</sub>	R <sub>3</sub>	R <sub>1</sub>	R <sub>2</sub>	R <sub>3</sub>	R <sub>1</sub>	R <sub>2</sub>	R <sub>3</sub>
P <sub>1</sub>	0	0	1	1	0	0	0	0	0
P <sub>2</sub>	1	0	1	0	1	0			
P <sub>3</sub>	0	1	0	0	0	1			

Request of none of the processes can be fulfilled based on the availability.  $\therefore$  Deadlock.



	Allocation		Request		Availability	
	R <sub>1</sub>	R <sub>2</sub>	R <sub>1</sub>	R <sub>2</sub>	R <sub>1</sub>	R <sub>2</sub>
✓ P <sub>1</sub>	0	1	1	0	0	0
✓ P <sub>2</sub>	1	0	0	0	1	0
✓ P <sub>3</sub>	1	0	0	1	0	0
	1	1			1	1

\* If cycle exists & each resource has only one instance, deadlock occurs for sure. But if it has multiple instances, deadlock may occur.



## Deadlock Handling

- i) Deadlock Prevention
- ii) Deadlock Avoidance
- iii) Deadlock detection & Recovery
- iv) Deadlock ignorance (ostrich method)

### 1) Deadlock Prevention

deadlock can be prevented by violating any of the 4 necessary conditions

#### → Removal of mutual exclusion

- shared resources must be used.
- can't be implemented

#### → Removal of hold and wait either hold or wait, not both Three ways:

wait x - A process must acquire all the necessary resources before execution starts (not implementable) - inefficient

hold x - If a process is holding some resources & requesting for additional resource, then it must release the acquired resources first (may lead to starvation)

hold x  
wait  
for limited  
time - A process holding some resources must wait for another resource for a limited time only after which it must release all the resources held by it.



## → Removal of No-Preemption for high priority processes - system processes

Two ways:

- If a process is holding some resources & requesting for another resource that can't be immediately allocated then all the acquired resources will be preempted.

- If a process is holding some resources & requests a resource,

If resource is available (allocated)

not available (resources held → allocated to some other process that is waiting) - preempted from the waiting process & allocated to it.

## → Removal of Circular Wait

To request resources

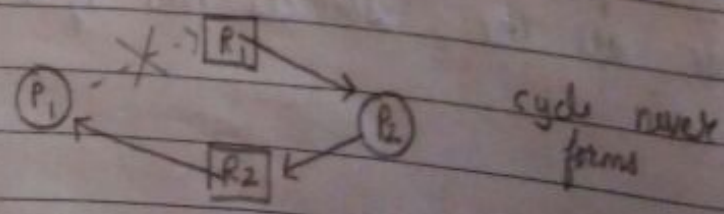
All the resources are associated with a no. To request resource  $R_j$ , a process must first release all  $R_i$  s.t.  $i \geq j$

A process requests & acquires resources in order

eg: Printer (1), CPU (5), Memory (6), CD drive (7)

$P_1$ : 1, 6  $\xrightarrow{\text{Request(5)}}$  1  $\xrightarrow{\text{(release 6)}}$  1, 5  $\xrightarrow{\text{(acquire 5)}}$  1, 5, 6  $\xrightarrow{\text{(acquire 6)}}$

(can acquire/request for 7)



Example:



prior info like  
 - no of res. allocated to a process  
 - max res required by a process  
 - order in which resources are required

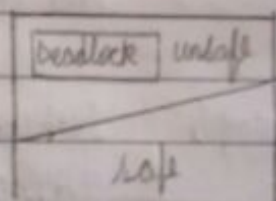
DL: \_\_\_\_\_ Pg: \_\_\_\_\_ B+

## 2) Deadlock Avoidance

- system maintains some database using which it can take decision whether to entertain a request or not s.t it remains in safe state
- system (kernel) analyse the database to determine whether granting a request can lead to deadlock in future
  - if not, then request is granted
  - otherwise keep pending until they can be granted (process may face long delay for obtaining a res.)

A safe seq. is possible

Safe state: if system can allocate all the resources requested by all the processes without entering into deadlock otherwise → unsafe state (cycle + not)



Initially system is in safe state A request

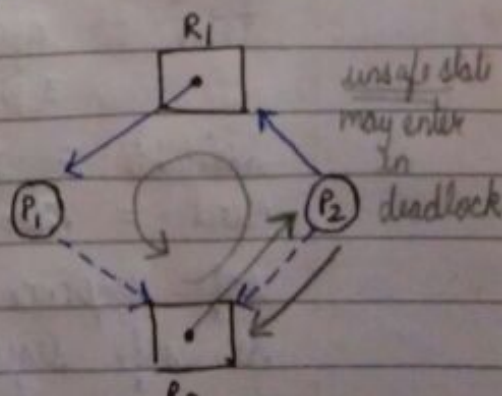
- Deadlock avoidance using RAG (used when all resources have only 1 instance)
  - $P_i \rightarrow R_j$  (request edge)
  - $R_j \rightarrow P_i$  (assignment edge)
  - $P_i \cdots \rightarrow R_j$   $P_i$  may request  $R_j$  in future (claim edge)

if granted

after completion

Resources must be claimed in advance i.e., in the beg. all edges in RAG are claim edges

Example: If  $P_i$  requests  $R_j$  then request edge can only be converted into assignment edge if it doesn't form a cycle. Here, request of  $P_2$  can't be granted.





for safe state, atleast one safe sequence must exist.

## Banker's Algorithm

used when resources have multiple instances.

It requires:

- Max. requirement of instances of each resource by each process. (2D array)
- Instances of each resource held by each process (2D array)
- Total instances of each resource available in the system (1D array)
- More instances required by each process (2D array)

Example: Snapshot at  $T_0$

	Allocation				Max				Available				Need			
	A	B	C	D	A	B	C	D	A	B	C	D	A	B	C	D
✓ $P_0$	0	0	1	2	0	0	1	2	1	5	2	0	0	0	0	0
✓ $P_1$	1	0	0	0	1	7	5	0	1	5	3	2	0	7	5	0
✓ $P_2$	1	3	5	4	2	3	5	6	2	8	8	6	1	0	0	2
✓ $P_3$	0	6	3	2	0	6	5	2	2	14	12	8	0	0	2	0
✓ $P_4$	0	0	1	4	0	6	5	6	2	14	12	12	0	6	4	2
	2	3	10	12					3	14	12	12				

Need matrix? Available res.?

Is system in safe state? If yes find safe seq.

Total : 

A	B	C	D
3	14	12	12

Available = Total - allocated

Need = Max - allocated

safe sequence:  $P_0 P_2 P_3 P_4 P_1$

As safe sequence exists  $\therefore$  It is in safe state

All requirements of all the processes are fulfilled  
 $\hookrightarrow$  (no cycle + int)  
 $\hookrightarrow$  no deadlock



Example :

Snapshot at time  $T_0$

	Allocation				Max				Available				Total			
	A	B	C	D	A	B	C	D	A	B	C	D	A	B	C	D
$P_0$	2	0	0	1	4	2	1	2	3	3	2	1	12	12	8	10
$P_1$	3	1	2	1	5	2	5	2								
$P_2$	2	1	0	3	2	3	1	6								
$P_3$	1	3	1	2	1	4	2	4								
$P_4$	1	4	3	2	3	6	6	5								
	9	9	6	9												

Available Resource instances = Total - Total (Allocation)  
 Need = Max - Allocation

	Allocation				Max				Available				Need			
	A	B	C	D	A	B	C	D	A	B	C	D	A	B	C	D
✓ $P_0$	2	0	0	1	4	2	1	2	3	3	2	1	2	2	1	1
✓ $P_1$	3	1	2	1	5	2	5	2	5	3	2	2	2	1	3	1
✓ $P_2$	2	1	0	3	2	3	1	6	6	6	3	4	0	2	1	3
✓ $P_3$	1	3	1	2	1	4	2	4	7	10	6	6	0	1	1	2
✓ $P_4$	1	4	3	2	3	6	6	5	10	11	8	7	2	2	3	3
									12	12	8	10				

safe sequence :  $P_0 \rightarrow P_3 \rightarrow P_4 \rightarrow P_1 \rightarrow P_2$

As safe sequence exists, system is in safe state



if request from  $P_1$  arrives for  $(1, 1, 0, 0)$ , Can this request be granted immediately?

Request must not exceed the need.

$$(1, 1, 0, 0) \leq (2, 1, 3, 1)$$

Check if Request  $\leq$  Available

$$(1, 1, 0, 0) \leq (3, 3, 2, 1) \Rightarrow \text{True}$$

$\therefore$  System pretends to grant the request & checks whether system remains in the safe state after granting the request.

	Allocation				Max				Available				Need			
	A	B	C	D	A	B	C	D	A	B	C	D	A	B	C	D
✓ $P_0$	2	0	0	1	4	2	1	2	2	2	2	1	2	2	1	1
✓ $P_1$	4	2	2	1	5	2	5	2	4	2	2	2	1	0	3	1
✓ $P_2$	2	1	0	3	2	3	1	6	5	5	3	4	0	2	1	3
✓ $P_3$	1	3	1	2	1	4	2	4	6	9	6	6	0	1	1	2
✓ $P_4$	1	4	3	2	3	6	6	5	10	11	8	7	2	2	3	3
									12	12	8	10				

safe sequence :  $P_0 \rightarrow P_3 \rightarrow P_4 \rightarrow P_1 \rightarrow P_2$

As safe sequence exists, system remains in safe state after granting the request.  
Therefore, request can be granted immediately.



If request from  $P_4$  arrives for  $(0, 0, 2, 0)$ , can it be granted immediately?

Request i.e.  $(0, 0, 2, 0) \leq$  Need i.e.  $(2, 2, 3, 3)$

Check if Request  $\leq$  Available

$(0, 0, 2, 0) \leq (3, 3, 2, 1) \Rightarrow \text{True}$

System pretends to grant the request:

	Allocation				Max				Available				Need			
	A	B	C	D	A	B	C	D	A	B	C	D	A	B	C	D
$P_0$	2	0	0	1	4	2	1	2	3	3	0	1	2	2	1	1
$P_1$	3	1	2	1	5	2	5	2					2	1	3	1
$P_2$	2	1	0	3	2	3	1	6					0	2	1	3
$P_3$	1	3	1	2	1	4	2	4					0	1	1	2
$P_4$	1	4	5	2	3	6	6	5					2	2	1	3

safe sequence: Since, need of none of the processes can be fulfilled with available resources, safe sequence doesn't exist  $\Rightarrow$  system is in unsafe state.

Therefore, request can't be granted immediately.



## Algorithm

Input: - Processes

- any 2 out of (Max, Need, Allocation matrix)
- any 1 out of (Available, Total)

### Safety Algorithm:

1. Let  $Work$  &  $Finish$  be vectors of length  $m$  &  $n$ .

$Work = Available$

$Finish[i] = false$  for  $i = 0, 1, \dots, n-1$   
 $n = \text{no. of processes}$

2. Find  $i$  such that:

$Finish[i] = false$

$Need_i \leq Work$

If no such  $i$  exists, go to step 4.

3.  $Work = Work + Allocation$

$Finish[i] = true$

go to step 2.

4. If  $finish[i] = true$  for all  $i$ , then system is in safe state  
else, it is in unsafe state.

Time Complexity:  $O(mn^2)$

$m = \text{no. of resources}$

$n = \text{no. of processes}$



## Resource - Request Algorithm for Process $P_i$ :

1. Request = request vector for process  $P_i$   
If  $\text{Request} \leq \text{Need}$ ,  
    go to step 2  
Otherwise,  
    raise error as process has exceeded its max claim.
2. If  $\text{Request} \leq \text{Available}$ ,  
    go to step 3  
Otherwise,  
     $P_i$  must wait as resources are not available.
3. Pretend to allocate requested resources to  $P_i$  by modifying:
  - $\text{Available} = \text{Available} - \text{Request};$
  - $\text{Allocation}_i = \text{Allocation}_i + \text{Request};$
  - $\text{Need}_i = \text{Need}_i - \text{Request};$
4. Find whether system is in safe state after allocation
  - If safe  $\Rightarrow$  resources are allocated to  $P_i$
  - If unsafe  $\Rightarrow P_i$  must wait & the old resource - allocation state is restored.



### 3) Deadlock Detection & Recovery

- Allow the system to enter into deadlocked state
- System checks for deadlock periodically
- Deadlock is detected using deadlock detection algorithms & recover the system using some recovery techniques

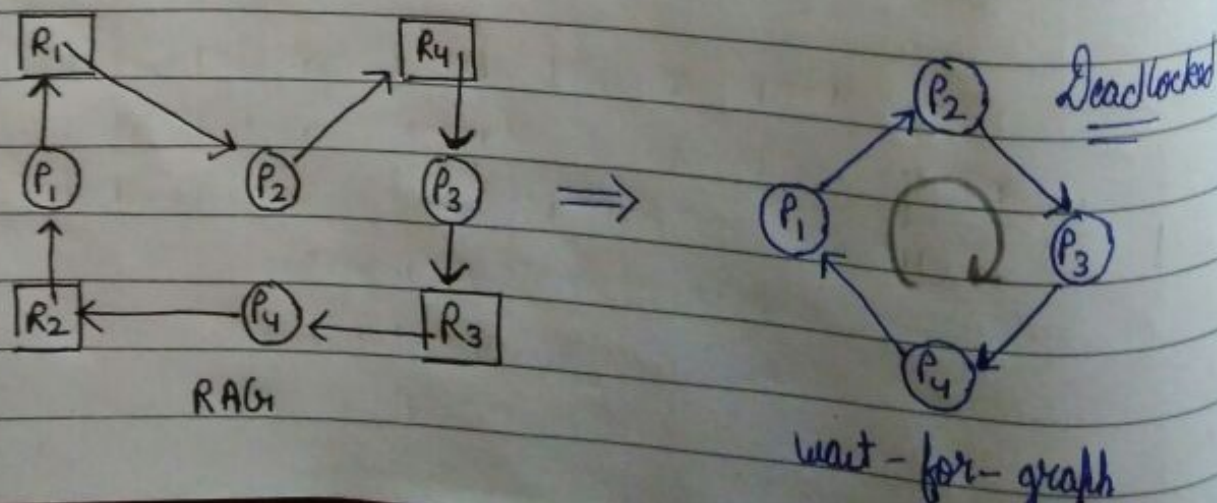
#### Deadlock detection algorithms

- i) Wait-for-graph (for single instances)
  - It is a RAG which consists of only processes as its nodes
  - set of
    - nodes - Processes ( $P_i$ )
    - edges -  $P_i \rightarrow P_j$   
( $P_i$  is waiting for the resource held by  $P_j$ )

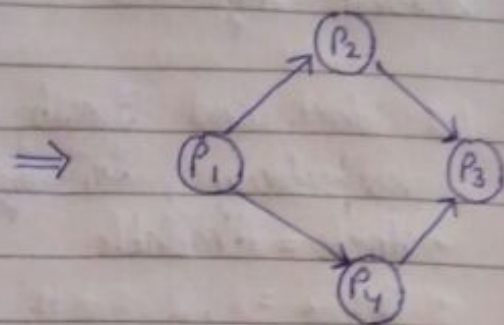
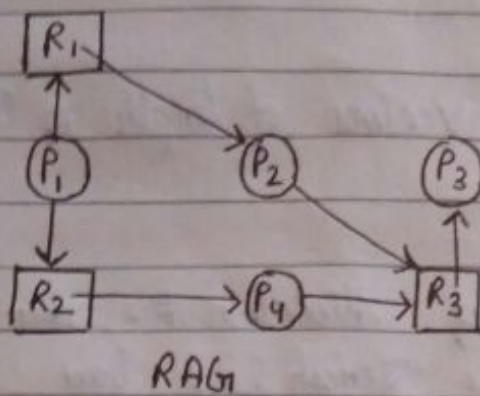
This algorithm checks for a cycle in the wait-for-graph

- If cycle is detected  $\Rightarrow$  system is in deadlock
- otherwise  $\Rightarrow$  system is not in deadlock

Example:







wait-for-graph  
(no cycle, no deadlock)

## 2) For multiple instances

It requires:

- Total no. of instances available for each resource in the system (1D array)
- Instances of each resource held by each process (2D array)
- Current request of instances of each resource by each process (2D array)

Example	Allocation			Request			Available		
	A	B	C	A	B	C	A	B	C
✓ P <sub>0</sub>	0	1	0	0	0	0	0	0	0
✓ P <sub>1</sub>	2	0	0	2	0	2	0	1	0
✓ P <sub>2</sub>	3	0	3	0	0	0	3	1	3
✓ P <sub>3</sub>	2	1	1	1	0	0	5	2	4
✓ P <sub>4</sub>	0	0	2	0	0	2	5	2	6
							7	2	6

safe sequence: P<sub>0</sub> → P<sub>2</sub> → P<sub>3</sub> → P<sub>4</sub> → P<sub>1</sub>  
system is not in deadlock

safe seq.: P<sub>0</sub> → further no request can be granted  
P<sub>1</sub>, P<sub>2</sub>, P<sub>3</sub>, P<sub>4</sub> are in deadlock



## Detection Algorithm

1) Let  $Work$  and  $Finish$  be vectors of length  $m$  &  $n$  respectively, Initialize

- $Work = Available$

- For  $i = 1, 2, \dots, n$ , if  $Allocation_i \neq 0$ , then  $Finish[i] = false$ , else  $Finish[i] = true$

2) Find an index  $i$  s.t. both

- $Finish[i] == false$

- $Request_i \leq Work$

If no such  $i$  exists, go to step 4

3)  $Work = Work + Allocation_i$

$Finish[i] = true$

go to step 2

4) If  $Finish[i] == false$  for some  $1 \leq i \leq n$ , system is in deadlock state &  $P_i$  is deadlocked

Time Complexity =  $O(mn^2)$



## Recovery Techniques

Recovery Techniques are used to recover a system from deadlock when detected by some deadlock detection algo.

Two approaches:

### i) Process Termination (Pessimistic Approach)

#### i) Abort all deadlocked processes.

- All the partially completed processes not in deadlock are aborted even if they have computed for a long time.
- This leads to high expenses.

#### ii) Abort one process at a time

- Abort one process at a time and run deadlock detection algorithm. If proc system is still in deadlocked state, abort another process and continue till deadlock is removed.
- Overhead of executing deadlock detection algo multiple times.
- Process to be aborted is decided on the basis of:
  - priority of processes.
  - time for which the process has computed.
  - time required for completion.
  - No. & type of resources utilised.
  - No. & type of resources needed.



Pg \_\_\_\_\_

2) Resource Preemption (optimistic approach)  
Pre-empt some resources from processes & give those resources to other processes until deadlock cycle is broken.  
Three issues:

i) Selecting a victim  
resources to be preempted must be decided in accordance with the cost minimization.

ii) Rollback

The process from which resources are preempted must be either

- rolled back completely & restarted again or
- rolled back to a safe state

iii) Starvation

If a process is picked as a victim multiple times, it may lead to starvation.

Therefore a process must be selected as a victim for only a finite no. of times.

4) Deadlock Ignorance

- It is used in the systems in which deadlock occurs rarely. deadlock is completely ignored and rather than spending resources on it, more frequently occurring issues are handled such as compiler errors, system crashes due to hardware failure, OS bugs.
- This approach is used in UNIX & Windows.
- If deadlock occurs, it is handled by rebooting.